

Why AI workloads are moving onto Kubernetes

The strongest adoption signal is about inference. The CNCF 2025 Annual Cloud Native Survey was released on 20 January 2026. It found that 66% of organisations hosting generative AI models use Kubernetes for some or all of their inference workloads, and that 82% of container users run Kubernetes in production, up from 66% in 2023^{1,2}. Official materials give only the combined 66% figure. We could not confirm a reported split into 'all' versus 'some' inference in any source^{1,3,4,44}.

Beneath that headline, adoption is shallow. Only 7% of organisations deploy AI models daily, 47% deploy them occasionally, and 44% do not yet run AI/ML workloads on Kubernetes at all^{1,3}. A CNCF blog from August 2026 cites 2025 platform-engineering research finding that 35% of platform teams still do not orchestrate AI workloads³⁸.

Table 1 Key figures from the CNCF 2025 Annual Cloud Native Survey

METRIC	VALUE	NOTE
Container users running Kubernetes in production	82%	Up from 66% in 2023
GenAI-hosting organisations using Kubernetes for some or all inference	66%	No official all/some split
Organisations not yet running AI/ML on Kubernetes	44%	Official announcement
Organisations deploying AI models daily	7%	47% deploy occasionally
Organisations that do not train models	Over half	Linux Foundation blog

Source: [1], [2], [3], [4], [44]. Data as of 2025 survey, published 2026-01-20. Secondary coverage reports a sample of 628 respondents. The official CNCF and Linux Foundation pages do not state the sample size^{1,4,45,46}.


NO TRAINING-ADOPTION FIGURE

No official CNCF or Linux Foundation page gives a share of organisations training models on Kubernetes. The 36% 'training' figure in the press release is a skills challenge, not model training^{1,3,4}. Public survey data covers inference far better than training^{1,3,47}.

Spending is shifting to inference

Gartner (August 2026) forecasts worldwide AI-optimised IaaS spending to grow 96% to about \$42B in 2026. Of that, \$23.3B (55%) is forecast for inference and \$19B for training. Gartner expects the inference share to reach 59% in 2027, in a market of about \$66B⁵. Deloitte measures compute rather than spending. It estimates inference at about one-third of AI compute in 2023 and about half in 2025, and projects about two-thirds in 2026⁴⁸.

Figure 1 Inference share of AI-optimised IaaS spending (Gartner forecast)
% of spending



Source: [5]. Data as of Forecast published 2026-08-10. Estimates, not measured spending.

The two forecasts measure different things and neither is a precise market share, but they point the same way. This suggests that the elastic, service-style workloads Kubernetes handles best will make up a growing share of AI infrastructure^{5,48}.